

```
In [1]: import pandas as pd

In [3]: teams = pd.read_csv("teams.csv")

In [3]: teams
Out[3]:
   team  country  year  events  athletes  age  height  weight  medals  prev_medals  prev_3_medals
0  AFG  Afghanistan  1964    8    8 22.0  161.0  64.2    0    0.0    0.0
1  AFG  Afghanistan  1968    5    5 23.2  170.2  70.0    0    0.0    0.0
2  AFG  Afghanistan  1972    8    8 29.0  168.3  63.8    0    0.0    0.0
3  AFG  Afghanistan  1980   11   11 23.6  168.4  63.2    0    0.0    0.0
4  AFG  Afghanistan  2004    5    5 18.6  170.8  64.8    0    0.0    0.0
...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...
2139 ZIM  Zimbabwe  2000   19   26 25.0  179.0  71.1    0    0.0    0.0
2140 ZIM  Zimbabwe  2004   11   14 25.1  177.8  70.5    3    0.0    0.0
2141 ZIM  Zimbabwe  2008   16   26 26.1  171.9  63.7    4    3.0    1.0
2142 ZIM  Zimbabwe  2012    8    9 27.3  174.4  65.2    0    4.0    2.3
2143 ZIM  Zimbabwe  2016   13   31 27.5  167.8  62.2    0    0.0    2.3

2144 rows x 11 columns

In [4]: teams = teams[["team", "country", "year", "athletes", "age", "prev_medals", "medals"]]

In [5]: teams
Out[5]:
   team  country  year  athletes  age  prev_medals  medals
0  AFG  Afghanistan  1964    8 22.0    0.0    0
1  AFG  Afghanistan  1968    5 23.2    0.0    0
2  AFG  Afghanistan  1972    8 29.0    0.0    0
3  AFG  Afghanistan  1980   11 23.6    0.0    0
4  AFG  Afghanistan  2004    5 18.6    0.0    0
...  ...  ...  ...  ...  ...  ...
2139 ZIM  Zimbabwe  2000   19 26 25.0    0.0    0
2140 ZIM  Zimbabwe  2004   11 14 25.1    0.0    3
2141 ZIM  Zimbabwe  2008   16 26 26.1    3.0    4
2142 ZIM  Zimbabwe  2012    8 9 27.3    4.0    0
2143 ZIM  Zimbabwe  2016   13 31 27.5    0.0    0

2144 rows x 7 columns

In [9]: correlation = teams["medals"].corr(teams["year"])
print(correlation)
-0.6216032363438134

In [11]: correlation = teams["medals"].corr(teams["athletes"])
print(correlation)
0.848574672847653

In [12]: correlation = teams["medals"].corr(teams["age"])
print(correlation)
0.6259858254295867

In [13]: correlation = teams["medals"].corr(teams["prev_medals"])
print(correlation)
0.6266222602328

In [14]: correlation = teams["medals"].corr(teams["medals"])
print(correlation)
1.0

In [16]: pip install seaborn

Collecting seaborn: you may need to restart the kernel to use updated packages.
Obtaining dependency information for seaborn from https://files.pythonhosted.org/packages/7b/65/83fcd7e9db836c170e8352bfc2d20f1d728197a16f883e709997a08e65e/seaborn-0.13.0-py3-none-any.whl.metadata
Downloading seaborn-0.13.0-py3-none-any.whl.metadata (5.3 kB)
Requirement already satisfied: numpy>=1.24.0,<=1.29 in c:\users\genet\appdata\local\programs\python\python32\lib\site-packages (from seaborn) (1.26.2)
Requirement already satisfied: pandas>=2.0 in c:\users\genet\appdata\local\programs\python\python32\lib\site-packages (from seaborn) (2.0.1.3)
Collecting matplotlib>=3.8.1,<=3.8.3 (from seaborn)
  Obtaining dependency information for matplotlib from https://files.pythonhosted.org/packages/2e/51/c77a14809976b9d6f4a40e31b0734f3955d6806c38ac5e708326b74415/matplotlib-3.8.2-cp312-cp312-win_amd64.whl.metadata
  Downloading matplotlib-3.8.2-cp312-cp312-win_amd64.whl.metadata (5.9 kB)
Obtaining dependency information for contourpy>=1.0.1 from https://files.pythonhosted.org/packages/b9/ae/a83308d40bf7a59228a6e08275f0a15c25e40d095b3383a36/contourpy-1.2.0-cp312-cp312-win_amd64.whl.metadata
  Downloading contourpy-1.2.0-cp312-cp312-win_amd64.whl.metadata (5.8 kB)
Collecting cycleor>=8.2.1, <9.0 (from matplotlib>=3.8.1,<=3.8.3->seaborn)
  Obtaining dependency information for cycleor from https://files.pythonhosted.org/packages/e7/06/c10818d9c309284af79471096926ef31630866e5114136c303d54a6e/cycleor-8.12.1-py3-none-any.whl.metadata
  Downloading cycleor-8.12.1-py3-none-any.whl.metadata (5.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib>=3.8.1,<=3.8.3->seaborn)
  Obtaining dependency information for fonttools from https://files.pythonhosted.org/packages/71/00/46562c0df644eaa2066d92630671a098f7f987f78c5186c41086/fonttools-4.46.1-cp312-cp312-win_amd64.whl.metadata
  Downloading fonttools-4.46.1-cp312-cp312-win_amd64.whl.metadata (158 kB)
.....
..... 0.67206 4 MB 7 eta .....
..... 0.67158 4 MB 92.8 KB/s eta 0:00:01
..... 0.67206 4 MB 7 eta .....
Collecting kiwisolver>=3.1 (from matplotlib>=3.8.1,<=3.8.3->seaborn)
  Obtaining dependency information for kiwisolver from https://files.pythonhosted.org/packages/62/50/2766560f44da82d17375785050c0f0784c40e629d5aa7a217071362/kiwisolver-1.4.5-cp312-cp312-win_amd64.whl.metadata
  Downloading kiwisolver-1.4.5-cp312-cp312-win_amd64.whl.metadata (5.6 kB)
Requirement already satisfied: packaging>=20.0 in c:\users\genet\appdata\local\programs\python\python32\lib\site-packages (from matplotlib>=3.8.1,<=3.8.3->seaborn) (23.2)
Collecting pillow>=8 (from matplotlib>=3.8.1,<=3.8.3->seaborn)
  Obtaining dependency information for pillow from https://files.pythonhosted.org/packages/32/64/9708651670897080c65032167608c29d9fcd0aabb4640b6906d5127f6/pillow-10.1.0-cp312-cp312-win_amd64.whl.metadata
  Downloading pillow-10.1.0-cp312-cp312-win_amd64.whl.metadata (9.6 kB)
Obtaining dependency information for pytyping>=2.3.1 (from matplotlib>=3.8.1,<=3.8.3->seaborn)
  Obtaining dependency information for pytyping from https://files.pythonhosted.org/packages/39/02/8480ed05fcc08f13b3aee4220212f0f0800b00e042132679a0247/pytyping-3.1.1-py3-none-any.whl.metadata
  Downloading pytyping-3.1.1-py3-none-any.whl.metadata (5.1 kB)
Requirement already satisfied: python-dateutil<=2.7 in c:\users\genet\appdata\local\programs\python\python32\lib\site-packages (from matplotlib>=3.8.1,<=3.8.3->seaborn) (2.8.2)
Requirement already satisfied: pyparsing>=2022.1 in c:\users\genet\appdata\local\programs\python\python32\lib\site-packages (from pandas>=1.2->seaborn) (2023.2)
Requirement already satisfied: s3==1.5 in c:\users\genet\appdata\local\programs\python\python32\lib\site-packages (from python-dateutil<=2.7->matplotlib>=3.8.1,<=3.8.3->seaborn) (1.16.8)
Downloading seaborn-0.13.0-py3-none-any.whl (294 kB)
.....
..... 0.67294 6 MB 7 eta .....
..... 0.67294 6 MB 5.9 MB/s eta 0:00:01
..... 0.67294 6 MB 4.5 MB/s eta 0:00:01
..... 0.67294 6 MB 3.1 MB/s eta 0:00:00
Downloading matplotlib-3.8.2-cp312-cp312-win_amd64.whl (7.6 MB)
.....
..... 0.67176 7 MB 7 eta .....
..... 0.67176 7 MB 8.1 MB/s eta 0:00:01
..... 0.67176 7 MB 11.3 MB/s eta 0:00:01
..... 0.67176 7 MB 17.1 MB/s eta 0:00:01
..... 0.67176 7 MB 25.3 MB/s eta 0:00:01
..... 0.67176 7 MB 23.3 MB/s eta 0:00:01
..... 0.67176 7 MB 25.9 MB/s eta 0:00:01
..... 0.67176 7 MB 24.5 MB/s eta 0:00:00
Download contourpy-1.2.0-cp312-cp312-win_amd64.whl (187 kB)
.....
..... 0.67187 7 MB 11.1 MB/s eta 0:00:00
Download cycleor-8.12.1-py3-none-any.whl (5.3 kB)
.....
Download fonttools-4.46.1-cp312-cp312-win_amd64.whl (21.3 MB)
.....
..... 0.67171 1 MB 7 eta .....
..... 0.67171 1 MB 33.8 MB/s eta 0:00:01
..... 0.67171 1 MB 27.4 MB/s eta 0:00:00
Download kiwisolver-1.4.5-cp312-cp312-win_amd64.whl (45 kB)
.....
..... 0.67161 1 MB 12.0 MB/s eta 0:00:00
Download pillow-10.1.0-cp312-cp312-win_amd64.whl (12.6 MB)
.....
..... 0.67161 1 MB 7 eta .....
..... 0.67161 1 MB 39.9 MB/s eta 0:00:01
..... 0.67161 1 MB 25.3 MB/s eta 0:00:00
Download pyparsing-3.1.1-py3-none-any.whl (102 kB)
.....
..... 0.67159 1 MB 7 eta 0:00:00
..... 0.67159 1 MB 17.1 MB/s eta 0:00:00
Installing collected packages: pytyping, pillow, kiwisolver, fonttools, cycleor, contourpy, matplotlib, seaborn
Successfully installed contourpy-1.2.0 cycleor-8.12.1 fonttools-4.46.1 kiwisolver-1.4.5 matplotlib-3.8.2 pillow-10.1.0 pytyping-3.1.1 seaborn-0.13.0
[notice] A new release of pip is available: 23.2.1 -> 23.3.1
[notice] To update, run 'python.exe -m pip install --upgrade pip'

In [17]: import seaborn as sns

In [24]: sns.plot(x="athletes", y="medals", data=teams, fit_reg=True, ci=None)

Out[24]:
<seaborn.axisgrid.FacetGrid at 0x1b0c417f0>

medals
athletes

In [26]: sns.lmplot(x="age", y="medals", data=teams, fit_reg=True, ci=None)

Out[26]:
<seaborn.axisgrid.FacetGrid at 0x1b0c417f0>

medals
age

In [27]: teams.plot.hist(y="medals")

Out[27]:
<Axes: xlabel='Frequency'>

Frequency
medals

In [28]: teams[teams.isnull().any(axis=1)]

Out[28]:
   team  country  year  athletes  age  prev_medals  medals
19  ALB  Albania  1992    9 25.3    NaN    0
26  ALG  Algeria  1964    7 26.0    NaN    0
39  AND  Andorra  1976    3 28.3    NaN    0
50  ANG  Angola  1980   17 17.4    NaN    0
59  ANT  Antigua and Barbuda  1976   17 23.2    NaN    0
...  ...  ...  ...  ...  ...  ...
2092 VIN  Saint Vincent and the Grenadines  1988    6 20.5    NaN    0
2103 YAR  North Yemen  1984    3 27.7    NaN    0
2105 YEM  Yemen  1992    8 19.6    NaN    0
2112 YMD  South Yemen  1988    5 23.6    NaN    0
2129 ZMA  Zambia  1984   15 21.7    NaN    0

130 rows x 7 columns

In [29]: teams=teams.dropna()

In [29]: teams
Out[29]:
   team  country  year  athletes  age  prev_medals  medals
0  AFG  Afghanistan  1964    8 22.0    0.0    0
1  AFG  Afghanistan  1968    5 23.2    0.0    0
2  AFG  Afghanistan  1972    8 29.0    0.0    0
3  AFG  Afghanistan  1980   11 23.6    0.0    0
4  AFG  Afghanistan  2004    5 18.6    0.0    0
...  ...  ...  ...  ...  ...  ...
2139 ZIM  Zimbabwe  2000   19 26 25.0    0.0    0
2140 ZIM  Zimbabwe  2004   11 14 25.1    0.0    3
2141 ZIM  Zimbabwe  2008   16 26 26.1    3.0    4
2142 ZIM  Zimbabwe  2012    8 9 27.3    4.0    0
2143 ZIM  Zimbabwe  2016   13 31 27.5    0.0    0

2014 rows x 7 columns

In [66]: train=teams[teams["year"]>=2012].copy()
newtest=teams[teams["year"]<2012].copy()

In [67]: train.shape
Out[67]:
(1689, 7)

In [76]: newtest.shape
Out[76]:
(405, 7)

In [64]: train
Out[64]:
   team  country  year  athletes  age  prev_medals  medals
0  AFG  Afghanistan  1964    8 22.0    0.0    0
1  AFG  Afghanistan  1968    5 23.2    0.0    0
2  AFG  Afghanistan  1972    8 29.0    0.0    0
3  AFG  Afghanistan  1980   11 23.6    0.0    0
4  AFG  Afghanistan  2004    5 18.6    0.0    0
...  ...  ...  ...  ...  ...  ...
2137 ZIM  Zimbabwe  1992   28 21.2    0.0    0
2138 ZIM  Zimbabwe  1996   21 23.8    0.0    0
2139 ZIM  Zimbabwe  2000   26 25.0    0.0    0
2140 ZIM  Zimbabwe  2004   14 25.1    0.0    3
2141 ZIM  Zimbabwe  2008   16 26.1    3.0    4

1609 rows x 7 columns

In [73]: from sklearn.linear_model import LinearRegression
reg = LinearRegression()

In [84]: predictors = ["athletes", "prev_medals"]
target = "medals"

In [85]: reg.fit(train[predictors], train["medals"])

Out[85]:
LinearRegression()

In [91]: predictions = reg.predict(newtest[predictors])

In [92]: newtest["predictions"] = predictions

In [92]: newtest
Out[92]:
   team  country  year  athletes  age  prev_medals  medals  predictions
6  AFG  Afghanistan  2012    6 24.8    1.0    1    -0.96121
7  AFG  Afghanistan  2016    3 24.7    1.0    0    -1.76333
24  ALB  Albania  2012   10 25.7    0.0    0    -1.42032
25  ALB  Albania  2016    6 23.7    0.0    0    -1.71747
37  ALG  Algeria  2012   39 24.8    2.0    1    2.10629
...  ...  ...  ...  ...  ...  ...
2111 YEM  Yemen  2016    3 19.3    0.0    0    -1.20268
2131 ZAM  Zambia  2012    7 22.6    0.0    0    -1.64043
2132 ZAM  Zambia  2016    7 24.1    0.0    0    -1.64043
2142 ZIM  Zimbabwe  2012    9 27.3    4.0    0    1.00767
2143 ZIM  Zimbabwe  2016   31 27.5    0.0    0    0.00748

405 rows x 8 columns

In [104]: newtest["newtest['predictions'] < 0, 'predictions'] = 0

In [105]: newtest["predictions"] = newtest["predictions"].round()

In [106]: newtest
Out[106]:
   team  country  year  athletes  age  prev_medals  medals  predictions
6  AFG  Afghanistan  2012    6 24.8    1.0    1    0.0
7  AFG  Afghanistan  2016    3 24.7    1.0    0    0.0
24  ALB  Albania  2012   10 25.7    0.0    0    0.0
25  ALB  Albania  2016    6 23.7    0.0    0    0.0
37  ALG  Algeria  2012   39 24.8    2.0    1    2.0
...  ...  ...  ...  ...  ...  ...
2111 YEM  Yemen  2016    3 19.3    0.0    0    0.0
2131 ZAM  Zambia  2012    7 22.6    0.0    0    0.0
2132 ZAM  Zambia  2016    7 24.1    0.0    0    0.0
2142 ZIM  Zimbabwe  2012    9 27.3    4.0    0    2.0
2143 ZIM  Zimbabwe  2016   31 27.5    0.0    0    0.0

405 rows x 8 columns

In [107]: from sklearn.metrics import mean_absolute_error
error = mean_absolute_error(newtest["medals"], newtest["predictions"])

In [108]: error
Out[108]:
3.288765432997655

In [109]: teams.describe()
Out[109]:
count    2814 0.00000
mean     18.00000
std      33.62728
min       0.00000
25%       0.00000
50%       0.00000
75%       0.00000
max      442.00000
Name: medals, dtype: float64

In [111]: newtest["team"] = "USA"

Out[111]:
   team  country  year  athletes  age  prev_medals  medals  predictions
2053 USA  United States  2012   689 26.7   317.0  248    285.0
2054 USA  United States  2016   719 28.4   246.0  264    236.0

In [112]: newtest["team"] = "IND"

Out[112]:
   team  country  year  athletes  age  prev_medals  medals  predictions
907 IND  India  2012    85 26.0    6.0    6    7.0
908 IND  India  2016   130 26.1    6.0    2   12.0

In [115]: errors = (newtest["medals"] - newtest["predictions"]).abs()

In [116]: errors
Out[116]:
6    1.0
7    0.0
24    0.0
25    0.0
37    1.0
...
2111    0.0
2131    0.0
2132    0.0
2142    2.0
2143    0.0
Length: 405, dtype: float64

In [117]: errors.groupby(newtest["team"]).mean()

In [118]: errors_by_team
Out[118]:
team
AFG    0.5
ALB    0.8
ALG    1.5
AND    0.0
ANG    0.0
...
VIN    0.0
YEM    0.0
ZAM    0.0
ZIM    1.0
Length: 284, dtype: float64

In [121]: medals_by_team = newtest["medals"].groupby(newtest["team"]).mean()
medals_by_team
Out[121]:
team
AFG    1.0
ALB    0.8
ALG    1.0
AND    NaN
ANG    NaN
...
VIN    NaN
YEM    NaN
ZAM    NaN
ZIM    1.0
Length: 162, dtype: float64

In [127]: import numpy as np
error_ratio = error_ratio(np.infinite(error_ratio))

In [128]: error_ratio
Out[128]:
team
AFG    1.000000
ALG    1.000000
AND    0.000000
ANG    0.000000
ARM    0.000000
AUS    0.007347
...
UZB    0.000000
VIE    0.000000
VIN    0.000000
VIE    1.000000
ZAM    0.000000
ZIM    1.000000
Length: 97, dtype: float64

In [129]: error_ratio.plot.hist()

Out[129]:
<Axes: xlabel='Frequency'>

Frequency
error_ratio.sort_values()

In [130]:
team
AFG    0.022472
CAN    0.048387
NZL    0.083462
RUS    0.023533
ITA    0.021429
...
POR    2.000000
EGY    2.400000
HKG    3.000000
LBR    3.333333
AUT    4.000000
Length: 97, dtype: float64

In [1]:
```